



A. The gate promotes models that are genuinely worse



B. Barely sensitive to q (at $\delta = -0.10$)



— $r = 0.25$
 — $r = 0.5$ (shipped default)
 — $r = 1.0$
 — $r = 2.0$
 - - fail-open path

Gate: accept iff $\text{macro-F1}_{\text{challenger}} > \text{macro-F1}_{\text{incumbent}} + 0.001$ (model_pipeline.py:447), both scored on the test split of a pool that is $1/(1+r)$ flagged incumbent errors and $r/(1+r)$ rows labelled by the incumbent itself. Monte-Carlo, 300 trials/point, pool $n=1500$, 3 classes, admin reporting rate $d=0.5$, challenger modelled as independent of the incumbent (conservative). Fail-open: model_pipeline.py:425-446.